SCOUT AI

AI Detector User Guide

Last Updated: 22 September 2025

SCOUTAI.NETLIFY.APP

Version: 1.1
Audience: First-time and returning users evaluating whether text is AI-generated.
Product: SCOUT AI - Advanced AI Content Detection System
Contact: sadyageraldm@gmail.com

How It Works

Understanding whether content was written by a human or by a large language model (LLM) matters for academic integrity, IP protection, compliance, and trust. This guide answers the most common questions about how our detector operates.

How is it possible to determine if something is AI-generated?

LLMs produce text by following statistical patterns learned from vast datasets. Compared to human writing, AI output often shows different distributions of phrasing, repetition, sentence structure, and “burstiness.”
Our detector analyzes those patterns at the document and span level. When sufficient text is provided, we assign a document-level likelihood and highlight spans that are most likely AI-generated with red indicators.

What’s the basic workflow?

1. **Input**: Paste text or upload a .txt/.docx file.
2. **Pre-checks**: We enforce **80–5,000 words** for reliable scoring.
3. **Analysis**: The backend computes a document score and span candidates.
4. **Presentation**:
    - Right panel: overall **% likely AI** + **Verdict** + actions (copy/download).
    - Left panel: your text with **red highlights** where AI-likelihood is highest.

What file types and lengths are supported?

- **File types:** .txt, .docx
- **Length: minimum 80 words**, **maximum 5,000 words** per submission
  (Shorter texts don't contain enough signal; extremely long texts increase noise and runtime.)

Does formatting change the outcome?

No. The detector evaluates **plain text**. Formatting (fonts, styles) is ignored after extraction.

Is any personal data retained?

Front-end text is sent to the backend **only** when you click **Detect**. Retention depends on your deployment configuration. For coursework submissions, follow your institution's data policy.

Understanding the Results

Where do I see the results?

- **Right panel** shows a large **percentage** (0–100%), the **Verdict**, and action buttons.
- **Left panel** shows your original text with **red highlights** for the most AI-likely spans.

- **Legend:** Red indicators for AI-generated content, green for human-generated content

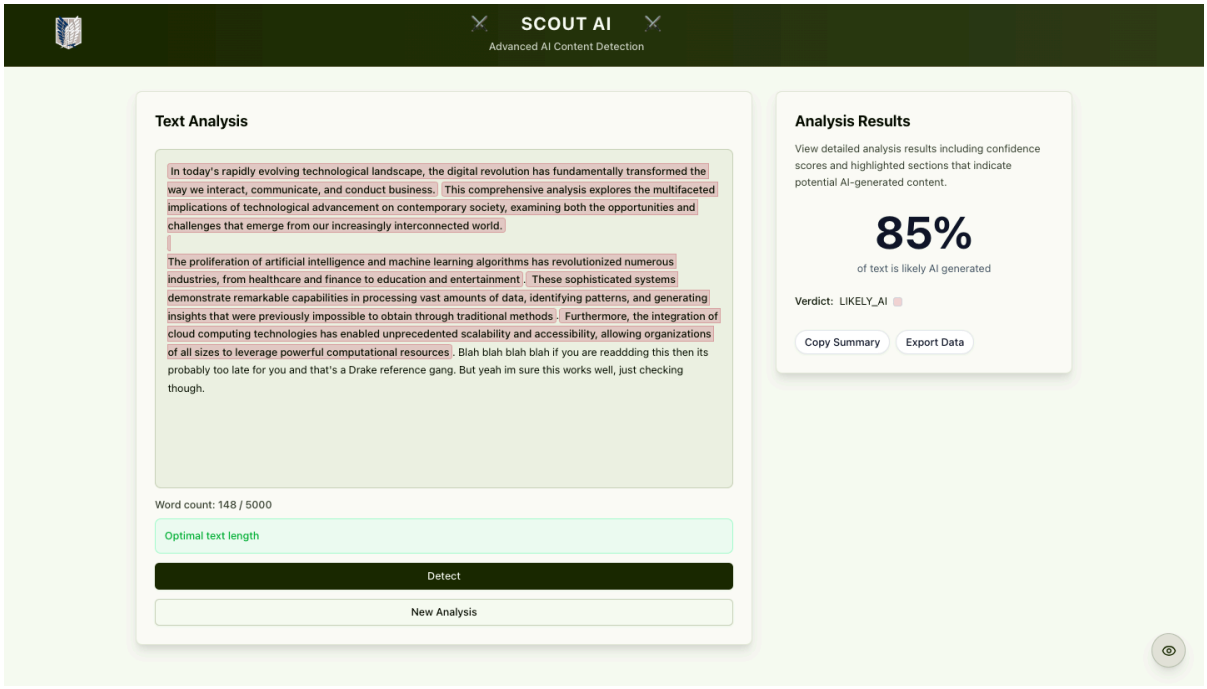

What does the percentage mean?

It's the overall **likelihood that the text contains AI-generated content**.
Example: **74%** means the detector finds strong AI-like patterns across a large portion of the document.

What verdicts can I get?

- **LIKELY AI** - strong evidence of AI-generated content.
- **LIKELY HUMAN** - insufficient evidence of AI generation.

The small colored square next to the verdict reinforces the status (e.g., red for **LIKELY AI**, green **LIKELY HUMAN**).

What do the highlights mean?

Red highlights mark **spans** whose patterns are most consistent with AI-generated text. Unmarked spans are comparatively more human-like. Highlights are **synchronized** with the exact text analyzed: if you edit the text, highlights disappear until you run **Detect** again.

Can you detect mixed text (human edited by AI or vice versa)?

Yes. Mixed authorship is common. You'll often see **partial highlighting** within paragraphs when only specific phrases look AI-generated.

Detection Capabilities & Limitations

What the detector does well

- **Expository & academic prose** with clear structure (paragraphs, topic sentences).
- **Span-level explanations** via highlights to support investigation.
- **Portable summary** and **JSON export** for audit trails.

Known limitations (be realistic and fair)

- **Short inputs** (<80 words) are unreliable — signals are too weak.
- **Highly stylized/creative writing** (poetry, lyrics) can be harder to classify.
- **Heavy paraphrasing or obfuscation** may reduce accuracy.
- **Non-English text**: accuracy depends on the underlying model and training; treat results with added caution.

> **Important:** AI detection should **inform**, not solely determine, high-stakes decisions. Combine detector output with context and human judgment.

FAQs

Why is there a minimum of 80 words?

Detection relies on statistical patterns that only emerge with enough text. Below 80 words, variance overwhelms signal, so we disable Detect and show how many more words you need.

Do .docx and .txt give different results?

No. We extract **plain text** from both; formatting is ignored. If a file can't be extracted (corrupt, password-protected), paste the content directly.

Will the detector flag grammar fixes from tools like spell-checkers?

Simple **spelling/grammar corrections** usually don't change authorship signals enough to flip a verdict. However, tools that **rewrite or paraphrase** with generative AI may increase AI-likelihood.

I edited after detection and my highlights disappeared—why?

Highlights match the **exact analyzed text**. Any edit invalidates the map. Click **Detect** again to re-align.

Can I trust a single percentage?

Treat it as **evidence**, not proof. For borderline cases, review highlighted spans, rerun on different sections, and consider external context (assignment instructions, writing history, etc.).

Can I export the results?

Yes. Click **Copy Summary** for a one-line summary, or **Export Data** for the full payload (verdict, score, highlights, timestamps, and the analyzed text).

What about privacy?

By default, text is sent to the backend only on **Detect**. Storage/retention depends on your deployment. For coursework, follow your institution's policy; for internal pilots, consult your admin.

Quick Start (Step-by-Step)

1. **Paste or Upload**
    - Paste directly, or click **Upload File** and choose .txt/.docx.
    - The file name appears in a badge; extracted text fills the editor.

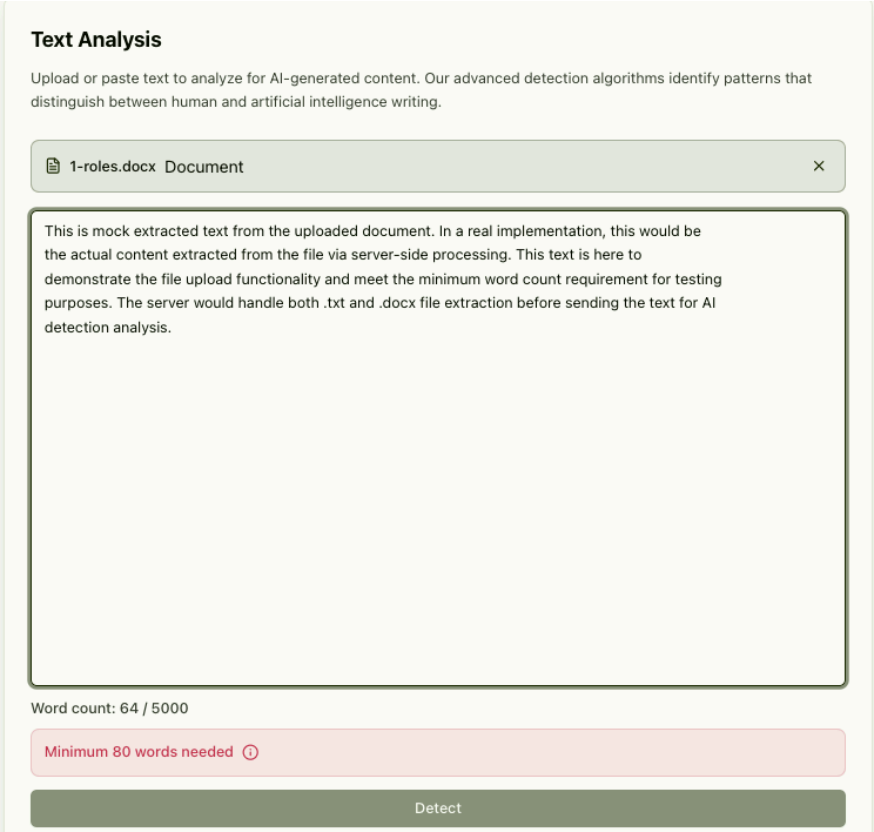

2. **Check Word Count**
    - Status under the editor shows **Minimum 80 words needed**, **Optimal text length**, or **Maximum 5,000 exceeded**.

- The **Detect** button enables only within range.

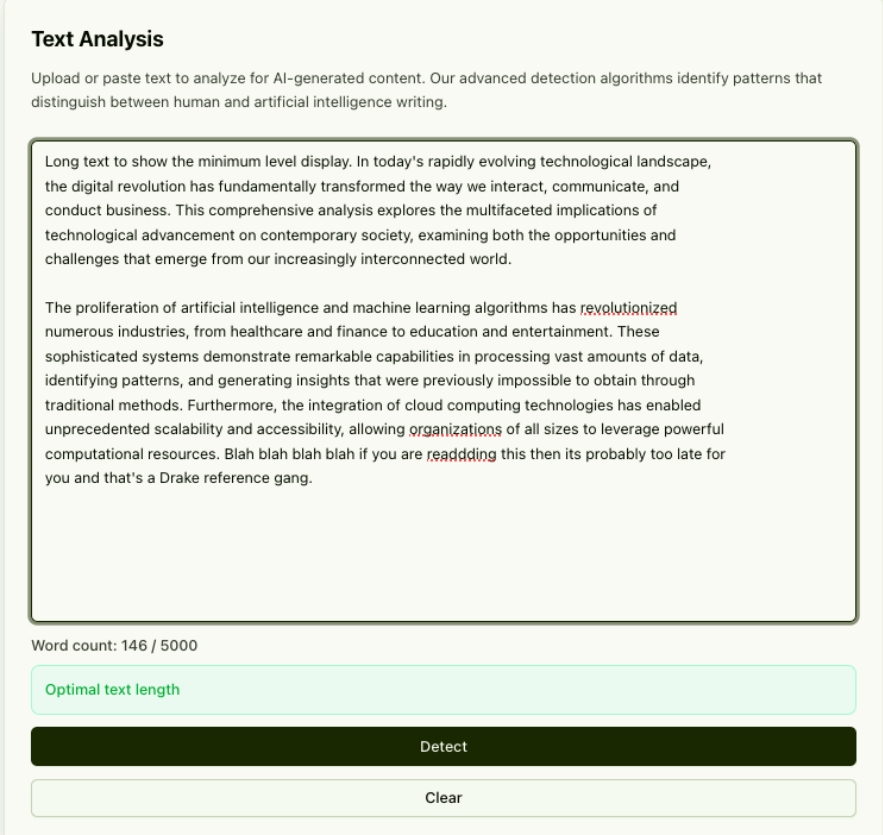

3. **Detect**
    - Click **Detect**. Right panel shows **Analyzing...**, then the percentage and verdict.
    - Left panel shows **red highlights** on AI-likely segments.

4. **Act on Results**
    - **Copy Summary** for quick sharing.
    - **Export Data** to archive or attach to reports.

5. **New Analysis**

- Click **New Analysis** (Clear) to reset the editor and results.

Tips for Better Assessments

- Use **complete paragraphs** instead of isolated sentences.
- If a document has multiple sections, analyze **each section** separately.
- Avoid pasting **tables/diagrams** — convert to prose for best results.
- For borderline cases, re-run after removing boilerplate (e.g., references, captions).

SCOUT AI - Advanced AI Content Detection System

"Wings of Freedom" - Detecting AI with Military Precision